

Resilient SADAD Payment Network — EMAM Framework Report

Author: Fuad Al Alawi

Course: Design Resilient National System

Date: November 26, 2025

GitHub: <https://github.com/FuadAlawi/Resilient-SADAD-System>

Executive Summary

This report presents a comprehensive resilient system design for Saudi Arabia's SADAD payment network, implementing the EMAM framework (Understand, Practice, Master, Excellence). The design addresses critical national infrastructure requirements through multi-layered resilience combining preventive security, fault tolerance, rapid recovery, and continuous operation under adverse conditions.

The SADAD network processes millions of daily transactions for utilities, government services, and e-commerce, making it a cornerstone of Saudi Arabia's digital economy and Vision 2030 objectives. Any disruption directly impacts citizens' ability to fulfill financial obligations and access essential services.

Key Deliverables: - Multi-Availability Zone architecture with automatic failover across 3 regions - Comprehensive chaos engineering suite testing 5 critical failure scenarios - Sub-15-minute RTO (Recovery Time Objective) for critical services - Infrastructure as Code using Terraform for rapid disaster recovery - Kubernetes-orchestrated services with horizontal auto-scaling - Prometheus-based observability with SLO-driven alerting

This design is grounded in Islamic principles of Amanah (stewardship), Adl (fairness), and No harm).

افهام (Understand) — Why Resilience > Security, Threat Analysis

Saudi Context and Critical Infrastructure

The Kingdom of Saudi Arabia's digital payment infrastructure represents a critical national asset that underpins both economic activity and Vision 2030's digital transformation objectives. The SADAD network, along with mada and SARIE systems, processes over 50 million transactions monthly, supporting:

- **Utility Payments:** Electricity, water, telecommunications bills for millions of households
- **Government Services:** Tax payments, visa fees, municipal services, traffic violations
- **E-Commerce:** Online retail, food delivery, digital services representing SAR billions in annual volume
- **Interbank Settlements:** Real-time payment settlements between financial institutions

Any extended outage creates cascading societal impacts: - Citizens unable to pay critical bills, risking service disconnection - Government revenue collection delays affecting public services - E-commerce merchants losing sales during peak periods (Ramadan, Eid) - Erosion of public trust in digital payment systems

Historical Context: The 2012 Shamoon malware attack on Saudi Aramco demonstrated the destructive potential of cyber threats targeting Saudi infrastructure. This attack erased data on over 30,000 computers, highlighting the need for resilience beyond traditional security measures.

Regulatory Framework: The Saudi Arabian Monetary Authority (SAMA) mandates: - Business Continuity Planning (BCP) with tested disaster recovery procedures - Regular resilience testing including chaos engineering scenarios - Incident reporting within 4 hours of service degradation - Annual third-party audit of critical systems - Minimum 99.9% availability for payment processing services

Why Resilience Exceeds Security

While security focuses on **preventing** incidents, resilience acknowledges that failures are inevitable and focuses on:

- 1. Continuing Operations Under Stress** - Security: "How do we prevent this attack?" - Resilience: "How do we maintain 80% service capacity if attacked?"

Example: During a DDoS attack, security measures (rate limiting, WAF) prevent some traffic, but resilience (auto-scaling, CDN, graceful degradation) ensures legitimate users can still transact.

- 2. Rapid Recovery Over Perfect Prevention** - RTO (Recovery Time Objective): 15 minutes for payment API - RPO (Recovery Point Objective): 1 minute for transaction ledger - MTTR (Mean Time To Recovery): Target < 30 minutes

A resilient system recovers from a compromise in minutes using Infrastructure as Code, whereas a security-only approach might take hours to investigate and rebuild.

- 3. Graceful Degradation** - Payment processing continues even if reporting dashboards fail - Read-only mode for ledger queries if write capacity is impaired - Queue transactions during network partition, replay when connectivity restored

Islamic Principles in Resilient Design

Amanah - Trust and Stewardship: - Citizens entrust their financial data and critical transactions to SADAD - Operators have a sacred duty to protect this trust through reliable service - Implementation: End-to-end encryption, audit trails, transparent incident reporting

Adl - Justice and Fairness: - All citizens deserve equal access to payment services regardless of geography - No group should be disproportionately affected by outages - Implementation: Multi-region deployment ensures service continuity if one region fails

No harm (No Harm): - Minimize societal harm through rapid containment and recovery - Clear communication during incidents to prevent panic or financial hardship - Implementation: Automated failover, pre-approved runbooks, public status page

Comprehensive Threat Model

Using the STRIDE methodology enhanced with operational failure scenarios:

Threat Actors: 1. **State-Sponsored APT Groups:** Sophisticated, well-resourced, targeting national infra 2. **Cybercriminal Organizations:** Financially motivated ransomware, data exfiltration 3. **Insider Threats:** Disgruntled employees, social engineering victims 4. **Supply Chain Compromises:** Malicious dependencies, compromised CI/CD pipelines 5. **Environmental/Operational:** Natural disasters, misconfigurations, human error

Critical Assets: | Asset Category | Examples | Resilience Measures | |-----|-----|-----||
Payment API | REST endpoints, authentication service | Multi-AZ, circuit breakers, rate limiting || Transaction Ledger | Database, audit logs | Replicated across 3 AZs, PITR backups || Settlement Services | Bank connectors, clearing house links | Redundant network paths, message queuing || Cryptographic Keys | KMS, HSM, TLS certificates | Key rotation, backup KMS in DR region || Customer PII | Names, IBANs, transaction history | Encryption at rest/transit, access logging |

Attack Paths and Mitigations: 1. **Credential Theft → Privilege Escalation** - Threat: Stolen admin credentials used to access production systems - Resilience: Break-glass access with MFA, session recording, automated

deprovisioning

2. Lateral Movement

- Threat: Compromised POD spreads to other services
- Resilience: Network policies, microsegmentation, namespace isolation

3. Container Breakout

- Threat: Escape from container to underlying node
- Resilience: Pod Security Standards, read-only file systems, AppArmor/SELinux

4. Supply Chain Attack

- Threat: Malicious dependency in application code
- Resilience: SBOM generation, signed images, admission controllers

Environmental Failures: - **AZ Failure:** AWS Availability Zone outage affecting 33% of infrastructure - **Regional Failure:** Entire AWS region (me-central-1) becomes unavailable - **Network Partition:** Split-brain scenario where services can't communicate - **DNS Outage:** Route53 or internal DNS resolution failures - **Data Center Fire:** Physical destruction of infrastructure (DR scenario)

عملية (Practice) — 4 Rs and Swiss Cheese Model

The Four Rs of Resilience

Our implementation operationalizes the 4 Rs framework across all system layers:

1. Robustness: Building Strength Into the System Application-Level Robustness:

- **Idempotent Operations:** All payment processing endpoints use idempotency keys to prevent duplicate charges if requests are retried

```
java @PostMapping("/api/v1/payments") public PaymentResponse processPayment(@RequestHeader("Idempotency-Key") String idempotencyKey, @RequestBody PaymentRequest request) { // Check if this exact request was already processed Optional<Payment> existing = paymentRepository.findByIdempotencyKey(idempotencyKey); if (existing.isPresent()) { return PaymentResponse.fromPayment(existing.get()); } // Process new payment atomically return paymentService.processWithKey(request, idempotencyKey); }
```

- **Circuit Breakers:** Prevent cascading failures when dependent services degrade
 - Open circuit after 5 consecutive failures
 - Half-open state after 30 seconds to test recovery
 - Fallback to cached data or degraded functionality

- **Rate Limiting:** Protect against abuse and ensure fair resource allocation
 - Per-user limits: 100 requests/minute for standard tier
 - Per-IP limits: 1000 requests/minute to prevent DDoS
 - Exponential backoff with jitter for retries

- **Input Validation:** Strong typing and schema validation at API boundaries

Infrastructure Robustness: - Kubernetes resource requests and limits prevent resource starvation - Pod Disruption Budgets (PDB) ensure minimum availability during updates - Topology spread constraints distribute pods across failure domains

2. Redundancy: Eliminating Single Points of Failure Geographic Redundancy:

- **Multi-AZ Deployment:** Services deployed across 3 Availability Zones in me-central-1 - Each AZ is a physically separate data center with

independent power/cooling - Minimum 2 pod replicas per AZ for critical services - Traffic automatically routes around failed AZs through Kubernetes Service

- **Cross-Region DR:** Warm standby in me-south-1 (UAE)
 - Database replica with 5-minute replication lag
 - Pre-deployed application infrastructure
 - DNS failover capability through Route53

Component Redundancy: | Layer | Component | Redundancy Strategy | |-----|-----|-----| | Load Balancing | AWS ALB | Multi-AZ with health checks | | Application | Payment Pods | Min 6 replicas (2 per AZ), HPA up to 12 | | Database | RDS PostgreSQL | Multi-AZ with automated failover | | Cache | Redis Cluster | 3-node cluster with replication | | Message Queue | Amazon SQS | Regionally redundant by design | | KMS | AWS KMS | Regional service with cross-region backup keys |

Network Redundancy: - Dual NAT Gateways per AZ for outbound connectivity - Multiple VPN tunnels for on-premise connectivity - BGP-based routing with automatic failover

3. Resourcefulness: Adapting to Changing Conditions Automated Response Playbooks:

1. High Error Rate Detected:

- Alert: ErrorRate > 1% for 5 minutes
- Action 1: Auto-scale pods +50%
- Action 2: Enable circuit breakers if not already active
- Action 3: Page on-call engineer
- Action 4: If error rate > 5%, initiate gradual rollback

2. Database Performance Degradation:

- Alert: DB latency p95 > 500ms
- Action 1: Scale read replicas
- Action 2: Enable query result caching
- Action 3: Throttle non-critical background jobs
- Action 4: If latency > 2s, enable read-only mode

Feature Flags for Graceful Degradation: - Real-time fraud checks can be disabled to reduce latency - Analytics/reporting features disabled during peak load - Transaction limits lowered during capacity constraints

Break-Glass Procedures: - Pre-approved emergency changes (no CAB review required) - Escalation matrix with mobile contacts for 24/7 response - Automated provisioning of emergency capacity

4. Rapidity: Speed of Detection and Recovery **Detection Speed** (Target: < 2 minutes): - Prometheus scrapes metrics every 15 seconds - AlertManager evaluates rules every 30 seconds - PagerDuty notification within 30 seconds of alert firing - Health check failures trigger immediate pod replacement

Response Speed (Target RTO: 15 minutes): - Automated horizontal scaling responds in < 60 seconds - Pod replacement (failed health checks) completes in < 90 seconds - AZ failover (Kubernetes native) completes in < 3 minutes - Region failover (manual DNS update) completes in < 15 minutes

Recovery Automation:

```
# Infrastructure recovery via Terraform
terraform apply -auto-approve -target=module.eks
```

```

# Application redeployment via Kubernetes
kubectl rollout restart deployment/payment-tier1
kubectl rollout status deployment/payment-tier1 --timeout=5m

# Database recovery from snapshot
aws rds restore-db-instance-from-db-snapshot \
--db-instance-identifier sadad-prod-recovered \
--db-snapshot-identifier sadad-backup-$(date +%Y%m%d)

```

Swiss Cheese Model: Defense in Depth

Each layer provides independent protection; an attack must penetrate ALL layers to cause harm:

Layer 1: Prevent (Stop attacks before they reach the system) - CI/CD pipeline scans for vulnerabilities (Snyk, Trivy) - Software Bill of Materials (SBOM) generation for all dependencies - Container image signing and verification (Cosign/Notary) - Least privilege IAM roles and service accounts - Network firewalls and security groups

Layer 2: Detect (Identify anomalies quickly) - Prometheus metrics: latency, error rates, saturation, traffic (RED/USE) - Distributed tracing (OpenTelemetry) for request flow visualization - Structured application logs aggregated to CloudWatch - SLO-based alerting (availability, latency, errors) - Blackbox probes from external vantage points

Layer 3: Contain (Limit blast radius) - Kubernetes Network Policies restrict pod-to-pod communication - Namespace isolation separates production from staging - KMS key scoping limits which services can decrypt sensitive data - Pod Security Standards prevent privilege escalation - Resource quotas prevent resource exhaustion

Layer 4: Recover (Restore service rapidly) - Multi-AZ deployment survives AZ failures automatically - Automated backups with Point-In-Time Recovery (PITR) - DR automation scripts with runbooks - Transaction replay from event log if ledger corrupted - Idempotency keys prevent duplicate processing during recovery

Layer 5: Learn (Continuous improvement) - Blameless postmortems after every incident - Chaos engineering drills quarterly - Automated runbook updates based on incident learnings - Security audits and penetration testing annually - Change Advisory Board (CAB) reviews high-risk changes

Practical Chaos Engineering Scenarios

We implement 5 chaos scenarios aligned with Saudi-specific threat landscape:

Scenario 1: Pod Kill (Medium Severity) **Objective:** Verify Kubernetes self-healing and service continuity

```

name: pod-kill-scenario
severity: medium
steps:
  - description: Kill one payment pod
    action: pod-kill
    parameters:
      selector: app=resilient-sadad-network
      mode: one
  - description: Verify service remains available
    action: http-check
    parameters:

```

```

url: http://payment/api/v1/healthz
expect: 200

```

Expected Outcome: Pod recreated in < 90s, no dropped requests

Scenario 2: Network Partition (High Severity) **Objective:** Test split-brain handling and consensus protocols

```

name: network-partition-scenario
severity: high
steps:
  - description: Partition AZ-1 from AZ-2/3
    action: network-partition
    parameters:
      source_zone: me-central-1a
      target_zones: [me-central-1b, me-central-1c]
      duration: 5m
  - description: Verify leader election
    action: verify-leader
    parameters:
      expect_new_leader: true
      timeout: 60s

```

Expected Outcome: New leader elected in minority partition, no data loss

Scenario 3: AZ Outage (Critical Severity) **Objective:** Validate multi-AZ resilience under complete AZ failure

```

name: az-outage-simulation
severity: critical
steps:
  - description: Simulate complete AZ failure
    action: drain-zone
    parameters:
      zone: me-central-1a
  - description: Monitor service availability
    action: availability-check
    parameters:
      min_success_rate: 0.95
      duration: 10m

```

Expected Outcome: Traffic shifts to AZ-2/3, 95%+ availability maintained

Scenario 4: DDoS During Eid (High Severity) **Objective:** Test auto-scaling and rate limiting under extreme load

```

name: ddos-eid-scenario
severity: high
steps:
  - description: Ramp traffic to 10x normal
    action: traffic-generator
    parameters:

```

```

    target_rps: 50000
    duration: 15m
- description: Verify auto-scaling response
  action: verify-hpa
  parameters:
    expect_min_replicas: 12
    expect_max_cpu: 70%

```

Expected Outcome: HPA scales to max replicas, latency < 1s p95

Scenario 5: Shamoon-Class Malware (Critical Severity) **Objective:** Test containment and recovery from destructive malware

```

name: shamoon-simulation
severity: critical
steps:
- description: Simulate disk wiper on node
  action: destroy-node-disks
  parameters:
    node: worker-node-3
    dry_run: true
- description: Verify pod evacuation
  action: verify-pod-migration
  parameters:
    expect_all_pods_rescheduled: true
    timeout: 5m
- description: Verify terraform can rebuild node
  action: terraform-plan
  parameters:
    target: aws_eks_node_group.general

```

Expected Outcome: Pods migrate off compromised node, node rebuilt from IaC

اتقان (Master) — Resilient Architecture & Critical Services

System Architecture Overview

The SADAD resilient payment network is built on a cloud-native architecture leveraging AWS managed services and Kubernetes for container orchestration. The design prioritizes:

1. **Horizontal scalability:** Add capacity by deploying more pods, not bigger instances
2. **Geographic distribution:** Multi-AZ within region + cross-region DR
3. **Automated recovery:** Self-healing infrastructure with minimal human intervention
4. **Observability:** Comprehensive metrics, logs, and traces for rapid troubleshooting

Application Layer

Payment Processing Service (Spring Boot) The core payment service is a stateless REST API built with Spring Boot 3.3.5 and Java 17:

Key Features: - **Idempotent Processing:** Prevents duplicate charges during retries - **Circuit Breakers:** Resilience4j library provides automatic circuit breaking - **Rate Limiting:** Bucket4j for token bucket algorithm implementation - **Health Checks:** Actuator endpoints for Kubernetes readiness/liveness probes - **Metrics Export:** Micrometer exports metrics in Prometheus format

Configuration (src/main/resources/ application.yml):

```
spring:
  application:
    name: resilient-sadad-network

  management:
    endpoints:
      web:
        exposure:
          include: health,prometheus,metrics
    metrics:
      export:
        prometheus:
          enabled: true
  health:
    livenessState:
      enabled: true
    readinessState:
      enabled: true

resilience4j:
  circuitbreaker:
    instances:
      paymentProcessing:
        registerHealthIndicator: true
        slidingWindowSize: 10
        permittedNumberOfCallsInHalfOpenState: 3
        failureRateThreshold: 50
        waitDurationInOpenState: 30s
```

Chaos Monkey Integration (application-chaos.yml):

```
spring:
  profiles: chaos

  chaos:
    monkey:
      enabled: true
      assaults:
        level: 5
        latencyActive: true
        latencyRangeStart: 500
        latencyRangeEnd: 2000
```

```

exceptionsActive: true
killApplicationActive: false

```

Ledger Writer Service Asynchronous service for writing payment records to the database:

- Uses Spring's `@Async` for non-blocking writes
- Implements outbox pattern for reliable message delivery
- Guarantees at-least-once delivery semantics
- Supports replay from transaction log if data corrupted

Settlement Connector Integrates with SARIE (Saudi RTGS) and bank APIs:

- Maintains persistent connections with retry logic
- Queues settlement requests during network outages
- Reconciles settlements daily with automated reports

Infrastructure Layer (AWS + Terraform)

VPC and Networking Our Terraform configuration provisions a production-grade network:

- 70% for 30 seconds
- Scale-down delay: 5 minutes to avoid flapping

Observability and SLOs

Prometheus Monitoring **ServiceMonitor for Metrics Collection** (`kubernetes/service-monitor.yaml`):

```

apiVersion: monitoring.coreos.com/v1
kind: ServiceMonitor
metadata:
  name: payment-metrics
  labels:
    app: resilient-sadad-network
spec:
  selector:
    matchLabels:
      app: resilient-sadad-network
  endpoints:
    - port: http
      path: /actuator/prometheus
      interval: 15s

```

PrometheusRule for SLO Alerts (`kubernetes/prometheus-rules.yaml`):

```

apiVersion: monitoring.coreos.com/v1
kind: PrometheusRule
metadata:
  name: payment-slo-alerts
spec:
  groups:
    - name: slo-alerts
      interval: 30s
      rules:
        - alert: HighErrorRate
          expr: |
            (

```

```

sum(rate(http_server_requests_seconds_count{status=~"5.."}[5m]))
/
sum(rate(http_server_requests_seconds_count[5m]))
) > 0.001
for: 5m
labels:
  severity: critical
annotations:
  summary: "Error rate above SLO threshold"
  description: "Error rate is {{ $value | humanizePercentage }}"

- alert: HighLatency
  expr: |
    histogram_quantile(0.95,
      sum(rate(http_server_requests_seconds_bucket[5m])) by (le)
    ) > 0.5
  for: 5m
  labels:
    severity: warning
  annotations:
    summary: "p95 latency above SLO threshold"
    description: "p95 latency is {{ $value }}s"

- alert: HighCPUUsage
  expr: |
    avg(
      rate(container_cpu_usage_seconds_total{pod=~"payment-*"}[5m])
    ) > 0.8
  for: 10m
  labels:
    severity: warning
  annotations:
    summary: "High CPU usage detected"
    description: "Average CPU usage is {{ $value | humanizePercentage }}"

```

SLO Targets: | Metric | Target | Measurement Window | |——|——| ——————| Availability | 99.9% | 30-day rolling | | Latency (p95) | < 500ms | 5-minute window | | Error Rate | < 0.1% | 5-minute window | | CPU Saturation | < 80% | 10-minute average |

Data Persistence and Backup

Database (RDS PostgreSQL)

- **Multi-AZ Deployment:** Automatic failover to standby in different AZ
- **Automated Backups:** Daily snapshots with 7-day retention
- **Point-In-Time Recovery:** Restore to any second within backup window

- **Read Replicas:** Offload analytics queries from primary

Backup Strategy

Data Type	Backup Method	Frequency	Retention	RPO
Transaction Ledger	RDS Automated Backup	Daily	7 days	1 minute
Configuration	Git + S3 Versioning	On commit	Indefinite	0 (immutable)
Application Logs	CloudWatch Logs	Real-time	90 days	0
Metrics	Prometheus TSDB	15s scrape	15 days	15 seconds

جیم (Excellence) — Innovation & Vision 2030

- **Innovation**

- Self-healing with policy-as-code to auto-quarantine compromised pods.
- Proactive capacity predictions using telemetry.
- Immutable infra with rapid rehydration from IaC.

- **Vision 2030 alignment**

- Enable fintech ecosystem and cashless society goals with high uptime and trust.
 - Data residency & compliance by design; transparency and public confidence.
 - Ethical stewardship grounded in Islamic principles of fairness and harm minimization.
-

Recovery Procedures (Runbooks)

See `docs/runbooks/recovery-procedures.md` for malware containment, DR failover, and key rotation playbooks, including RTO/RPO targets and verification steps.

Chaos, Monitoring, and SLOs

- Run app with Chaos Monkey profile (local):
 - `mvn spring-boot:run -Dspring-boot.run.profiles=chaos`
 - `scripts/chaos-monkey-demo.sh`
 - Kubernetes monitoring:
 - Apply `kubernetes/service-monitor.yaml` and `kubernetes/prometheus-rules.yaml` with `kube-prometheus-stack`.
 - Scenarios:
 - See `chaos-tests/` YAMLs for pod kill, partition, zone outage, DDoS, and Shamoon simulation.
-

Verification Results

Application Build & Tests

Test Phase	Status	Details
Maven Build	<input type="checkbox"/> SUCCESS	Compiled 3 source files
Resources	<input type="checkbox"/> SUCCESS	Copied 2 resources to target/classes
Unit Tests	<input type="checkbox"/> SUCCESS	All tests passed
Build Time	<input type="checkbox"/> 1.045s	Fast build cycle
Date	<input type="checkbox"/> Verified	2025-11-26T20:03:15+03:00

Command: mvn clean test

Chaos Test Scenarios

Scenario	File	Severity	Status	Description
AZ Outage	az-outage-simulation.yml	<input type="checkbox"/> Critical	<input type="checkbox"/> Ready	Simulates complete Availability Zone failure
DDoS (Eid)	ddos-eid-scenarioHigh.yml	<input type="checkbox"/> High	<input type="checkbox"/> Ready	High-traffic scenario during Eid period
Network Partition	network-partitionscenario.yml	<input type="checkbox"/> Medium	<input type="checkbox"/> Ready	Simulates network split between zones
Pod Kill	pod-kill-scenarioMedium.yml	<input type="checkbox"/> Medium	<input type="checkbox"/> Ready	Kills payment pod, verifies availability
Shamoon Malware	shamoon-simulation.yml	<input type="checkbox"/> Critical	<input type="checkbox"/> Ready	Disk-wipe simulation with recovery

Command: bash chaos-tests/run-all-tests.sh

Total Scenarios: 5 (2 Critical, 2 High, 1 Medium)

Infrastructure Components

Component	Technology	Configuration	Resilience Features
VPC	AWS VPC	10.0.0.0/16	3 Availability Zones
Network	Public/Private Subnets	3 public + 3 private	Multi-AZ isolation
NAT	NAT Gateway	One per AZ	High availability
Compute	EKS 1.27	Auto Scaling	Min: 3, Max: 6 nodes
Nodes	t3.medium	ON_DEMAND	Spread across AZs
Monitoring	Prometheus	ServiceMonitor	SLO-based alerts
Orchestration	Kubernetes	HPA + PDB	Self-healing

Terraform Files: main.tf, variables.tf, outputs.tf, versions.tf

Submission

- Repository: include this EMAM report, runbooks, chaos scenarios, Terraform, and Kubernetes manifests.
- PDF: export this document via `scripts/build-pdf.sh` (requires pandoc) or Print-to-PDF.
- Checklist:
 - EMAM report covers Saudi context and Islamic principles.
 - 4 Rs and Swiss Cheese applied with practical tests.
 - Architecture, SLOs, and recovery procedures included.
 - Chaos tests runnable (stubs or integrated) and monitored.